

Db2 for z/OS Modern Application: Watson, ML, REST, Spark and more

Jane Man

IBM

Session code: G12

Wednesday, 04.10.2017 13:40-14:40

Platform: Db2 for z/OS



Objectives

- To introduce how to build a cognitive Db2 solution using Watson
- To illustrate how to do Machine Learning on Db2 data without a GUI
- To discuss how to do Spark analytics
- To show how to use Db2 as REST provider using Db2 native REST service
- To demonstrate how to use Db2 as REST consumer using JSON feature in Db2

Please note

- IBM's statements regarding its plans, directions, and intent are subject to change or withdrawal without notice and at IBM's sole discretion.
- Information regarding potential future products is intended to outline our general product direction and it should not be relied on in making a purchasing decision.
- The information mentioned regarding potential future products is not a commitment, promise, or legal obligation to deliver any material, code or functionality. Information about potential future products may not be incorporated into any contract.
- The development, release, and timing of any future features or functionality described for our products remains at our sole discretion.
- Performance is based on measurements and projections using standard IBM benchmarks in a controlled environment. The actual throughput or performance that any user will experience will vary depending upon many factors, including considerations such as the amount of multiprogramming in the user's job stream, the I/O configuration, the storage configuration, and the workload processed. Therefore, no assurance can be given that an individual user will achieve results similar to those stated here.

Agenda

- Db2 as REST provider using Db2 native REST service
- Db2 as REST consumer using JSON feature in Db2
- Cognitive Db2 solution using Waston
- Perform Machine Learning on Db2 data without a GUI
- Spark analytics

Agenda

- Db2 as REST provider using Db2 native REST service
- Db2 as REST consumer using JSON feature in Db2
- Cognitive Db2 solution using Waston
- Perform Machine Learning on Db2 data without a GUI
- Spark analytics

540 million
RESTful
requests per
hour

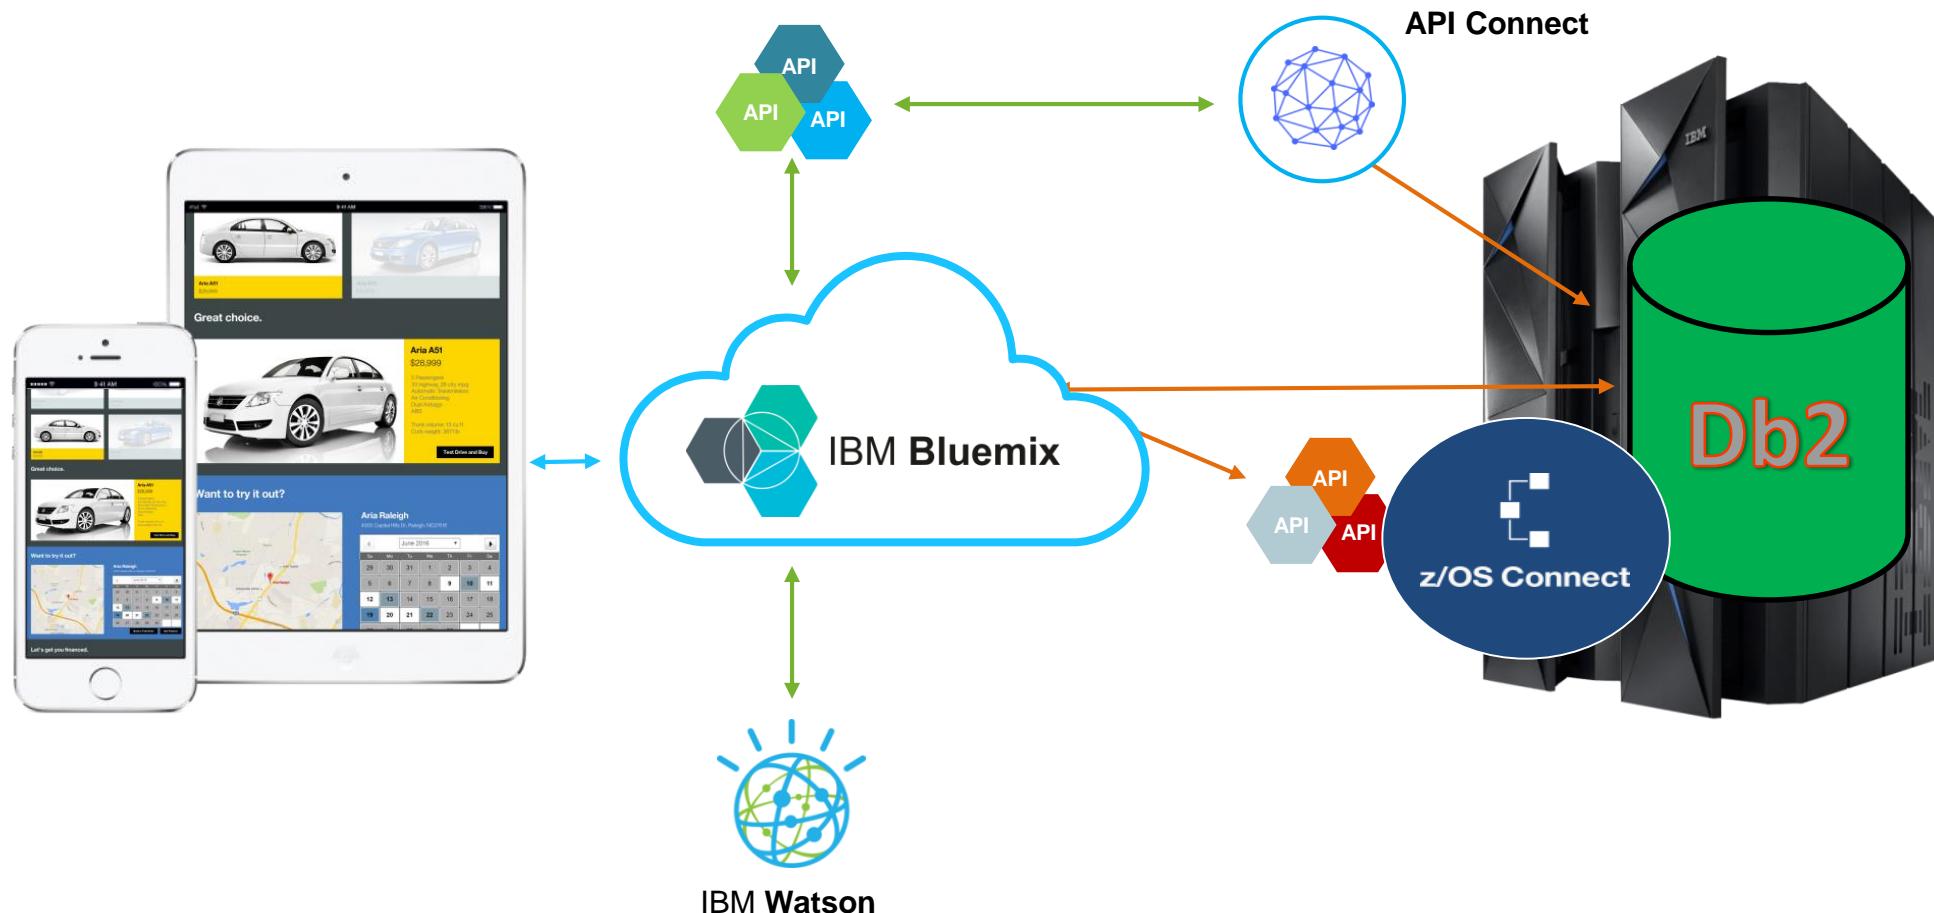
```
POST https://<host>:<port>/services/simpleSelect1  
{ "LOCATION": "<location>" }
```

Db2 is now a RESTful service provider (V11 PI66828)

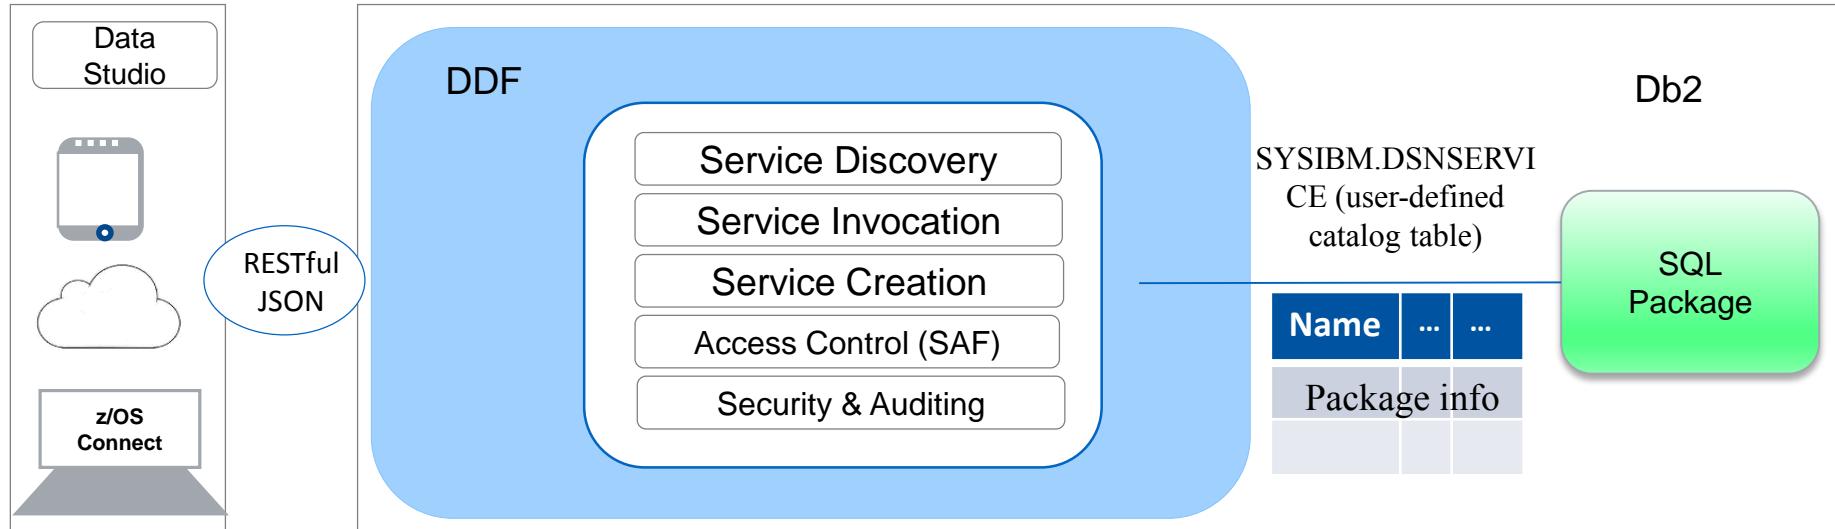
Enabling new business value for your enterprise data

Modernizing using the power of SQL

Unleashing Db2 data for the API Economy



Db2 as a Service Provider – Solution Overview



Db2 User Created Service – A single SQL statement (INSERT, UPDATE, DELETE, SELECT, CALL) bound into a package.

Db2 Management Services – Create, discover and invoke a DB2 user created service using a REST API.

DDF Interface – Processes HTTP requests and replies, performs REST/JSON message formatting invokes as a static SQL statement

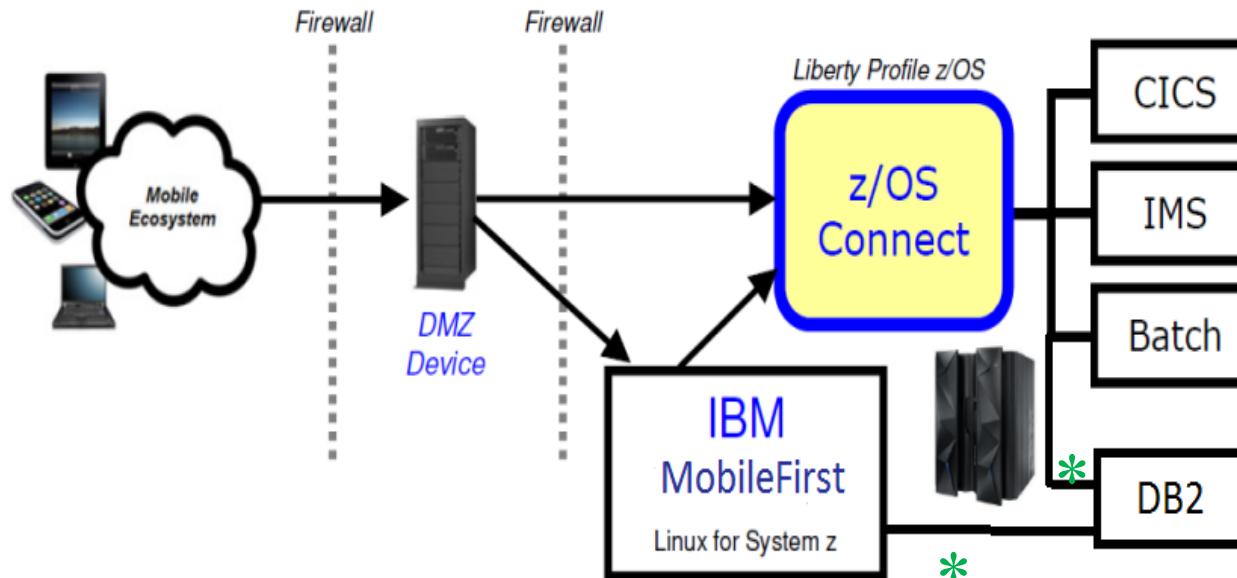
Availability – Utilizes existing DDF capability including thread pooling, profiling and enclave classification.

Leverages existing package lifecycle – Services are in the catalog tables and saved in the directory.

Security, accounting, statistics, and auditing – Uses existing DDF infrastructure. Supporting classification, profiling, and thread pooling

z/OS Connect EE Integration – Uses the z/OS Connect EE REST Client to allow exporting of DB2 svcs into z/OS Connect rest client

Simplified End-to-End Architecture for Mobile and Cloud Application invoking z Services using APIs



JDBC or **Db2**
Native REST

Create a Db2 REST service

1. Customize and Run DSNTIJRS to create Db2 RESTful service database and table
2. Authorize user to access services
Access to Db2 REST is protected by a SAF DSNR class <ssid>.REST profile
3. Create a Db2 REST service

```
//RESTDNSR JOB CLASS=A,MSGCLASS=H,REGION=0M,
//    USER=RACF000,PASSWORD=PASSWORD
//*****
//** RACF COMMANDS TO ADD DSNR CLASS PROFILES FOR REST
//*****
/STEP1 EXEC PGM=IKJEFT01
//SYSRACF DD DSN=SYS1.RACF,DISP=SHR
//SYSPRINT DD SYSOUT=*
//SYSTSPRT DD SYSOUT=*
//SYSTSIN DD *
RDEFINE DSNR (DB2A.REST) UACC(READ)
```

POST ▾ <http://dtec731.vmec.svl.ibm.com:446/services/DB2ServiceManager>

Headers

The headers specified below will be added to any sent by the browser or specified elsewhere.

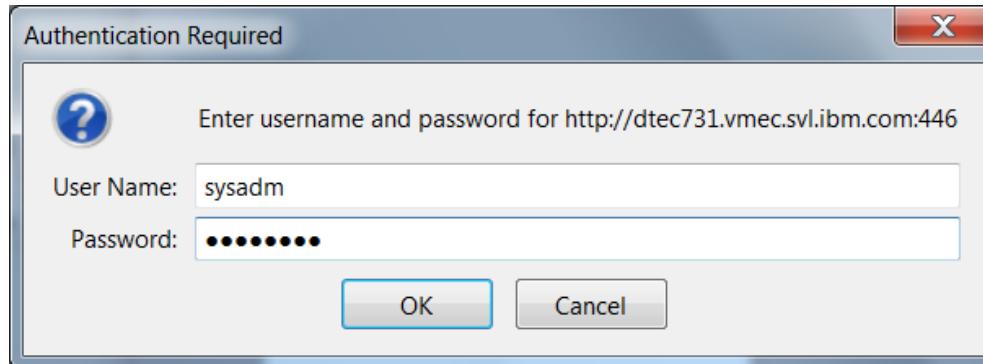
Name	Value
Accept	application/json
Content-Type	application/json
Accept-Charset	UTF-8

Enter the data and its corresponding MIME type below.

application/json

```
{
  "requestType": "createService",
  "sqlStmt": "SELECT SUBSTR(STRING,1,60) as STRING from
$YSIBM.SYSXMLSTRINGS WHERE STRINGID= ?",
  "serviceName": "selectSYSXMLStrings",
  "description": "Select string from XML string id"
}
```

Backend details : create a Db2 REST service



Output of creating Db2
REST service

```
{  
  "StatusCode": 201,  
  "StatusDescription": "DB2 Rest Service selectSYSXMLStrings was created  
successfully.",  
  "URL":  
  "http://dtec731.vmec.svl.ibm.com:446/services/SYSIBMSERVICE/selectSYSXMLStrings"  
}
```

Backend details : Discover all Db2 REST service (optional)

GET <http://dtec731.vmec.svl.ibm.com:446/services>

```
{  
    "DB2Services": [ {  
        "ServiceName": "DB2ServiceDiscover",  
        "ServiceCollectionID": null,  
        "ServiceDescription": "DB2 service to list all available services.",  
        "ServiceProvider": "db2service-1.0",  
        "ServiceURL": "http://dtec731.vmec.svl.ibm.com:446/services/DB2ServiceDiscover"  
    }, {  
        "ServiceName": "DB2ServiceManager",  
        "ServiceCollectionID": null,  
        "ServiceDescription": "DB2 service to create, drop, or alter a user defined  
service.",  
        "ServiceProvider": "db2service-1.0",  
        "ServiceURL": "http://dtec731.vmec.svl.ibm.com:446/services/DB2ServiceManager"  
    }, {  
        "ServiceName": "selectSYSXMLStrings",  
        "ServiceCollectionID": "SYSIBMSERVICE",  
        "ServiceDescription": "Select string from XML string id",  
        "ServiceProvider": "db2service-1.0",  
        "ServiceURL":  
        "http://dtec731.vmec.svl.ibm.com:446/services//SYSIBMSERVICE/selectSYSXMLStrings"  
    } ]  
}
```

Backend details : create a Db2 REST svs (cont'd)

4. Get details for the service

```
{
  "StatusCode": 201,
  "StatusDescription": "DB2 Rest Service selectSYSXMLStrings was created successfully.",
  "URL":
  "http://dtec731.vmec.svl.ibm.com:446/services/SYSIBMSERVICE/selectSYSXMLStrings"
}
```

Output of creating
Db2 REST service

GET <http://dtec731.vmec.svl.ibm.com:446/services/SYSIBMSERVICE/selectSYSXMLStrings>

```
{"selectSYSXMLStrings": {
    "serviceName": "selectSYSXMLStrings",
    "serviceCollectionID": "SYSIBMSERVICE",
    ...
    "serviceStatus": "started",
    "RequestSchema": {...,
        "properties": {
            "P1": {
                "type": ["null", "integer"],
                "multipleOf": 1,
                "minimum": -2147483648,
                "maximum": 2147483647,
                "description": "Nullable INTEGER"
            }
        },
        "required": ["P1"], ...
    }
}
```

INPUT

Backend details : create a Db2 REST service (cont'd)

Details on `selectSYSXMLStrings` service

```
"ResponseSchema": {  
    "$schema": "http://json-schema.org/draft-04/schema#",  
    "type": "object",  
    "properties": {  
         "ResultSet Output": {  
            "type": "array",  
            "items": {  
                "type": "object",  
                "properties": {  
                    "STRING": {  
                        "type": "string",  
                        "description": "CHAR(60)"  
                    }  
                },  
                "required": ["STRING"],...  
            }  
        }  
    }  
}
```

OUTPUT

Backend details : create a Db2 REST service (cont'd)

5. Invoke a service

POST <http://dtec731.vmec.svl.ibm.com:446/services/SYSIBMSERVICE/selectSYSXMLStrings>

- Headers

The headers specified below will be added to any sent by the browser or specified elsewhere.

Name	Value
Accept	application/json
Content-Type	application/json
Accept-Charset	UTF-8
Name	Value

- Data

Select one of the options below to include data with the request.

Custom

Enter the data and its corresponding content type:

application/json

```
{
  "P1": 1006
}
```

Why "P1"?

OUTPUT

```
{
  "ResultSet Output": [ {
    "STRING": "space" } ],
  "StatusCode": 200,
  "StatusDescription":
  "Execution Successful"
}
```

Agenda

- Db2 as REST provider using Db2 native REST service
- **Db2 as REST consumer using JSON feature in Db2**
- Cognitive Db2 solution using Waston
- Perform Machine Learning on Db2 data without a GUI
- Spark analytics

JSON_VAL Built-in function

To extract and retrieve JSON data into SQL data types from BSON.

```
>>-JSON_VAL(--json-value--,--search-string--,--result-type--)-----><
```

The `JSON_VAL` function returns an element of a JSON document identified by the JSON field name specified in *search-string*. The value of the JSON element is returned in the data type and length specified in *result-type*.

Example:

```
JSON_VAL(DATA,'PO.customer@cid', 'i:na')
```

<i>Result-type</i>	Function return type / length
'n'	DECFLOAT(34)
'i'	INTEGER
T	BIGINT
'f'	DOUBLE
'd'	DATE
'ts'	TIMESTAMP
't'	TIME
's:n'	VARCHAR (n)
'b:n'	VARCHAR(n) FOR BIT DATA
'u'	INTEGER / 4

PI39003 remove the requirement that 1st parameter has to be a BLOB column

SYSTOOLS.JSON_TABLE

This function returns array of elements in JSON data.

Example:

```
SELECT X.* FROM JSONPO,
TABLE(SYSTOOLS.JSON_TABLE(DATA,
'PO.items.item', 's:200')) X
```

Output:

TYPE	VALUE
3	{@partNum:"872-AA",productName:"Lawnmower",quantity:1,USPrice:149.990000,shipDate:"2014-11-20"}
3	{@partNum:"945-ZG",productName:"Sapphire Bracelet",quantity:2,USPrice:178.990000,comment:"Not shipped"}

```
{"PO": {"@id": 101,
        "@orderDate": "2014-11-18",
        "customer": {"@cid": 999},
        "items": [
            {"item": [{"@partNum": "872-AA",
                      "productName": "Lawnmower",
                      "quantity": 1,
                      "USPrice": 149.99,
                      "shipDate": "2014-11-20"
                    },
                    {"@partNum": "945-ZG",
                      "productName": "Sapphire Bracelet",
                      "quantity": 2,
                      "USPrice": 178.99,
                      "comment": "Not shipped"
                    }
                ]
        }
    }
```

JSON_TABLE + JSON_VAL : convert JSON to Relational

Example:

SELECT

```
JSON_VAL(SYSTOOLS.JSON2BSON(X.VALUE), '@partNum', 's:10') as "@partNum",
JSON_VAL(SYSTOOLS.JSON2BSON(X.VALUE), 'productName', 's:20') as "productName",
JSON_VAL(SYSTOOLS.JSON2BSON(X.VALUE), 'quantity', 'i') as "quantity",
JSON_VAL(SYSTOOLS.JSON2BSON(X.VALUE), 'USPrice', 'f') as "USPrice",
JSON_VAL(SYSTOOLS.JSON2BSON(X.VALUE), 'shipDate', 's:20') as "shipDate"
```

FROM JSONPO,

```
TABLE(SYSTOOLS.JSON_TABLE(DATA, 'PO.items.item', 's:200')) X
```



VALUE

```
{@partNum:"872-AA",productName:"Lawnmower",quantity:1,USPrice:149.990000,shipDate:"2014-11-20"}
{@partNum:"945-ZG",productName:"Sapphire Bracelet",quantity:2,USPrice:178.990000,comment:"Not shipped"}
```

Output:

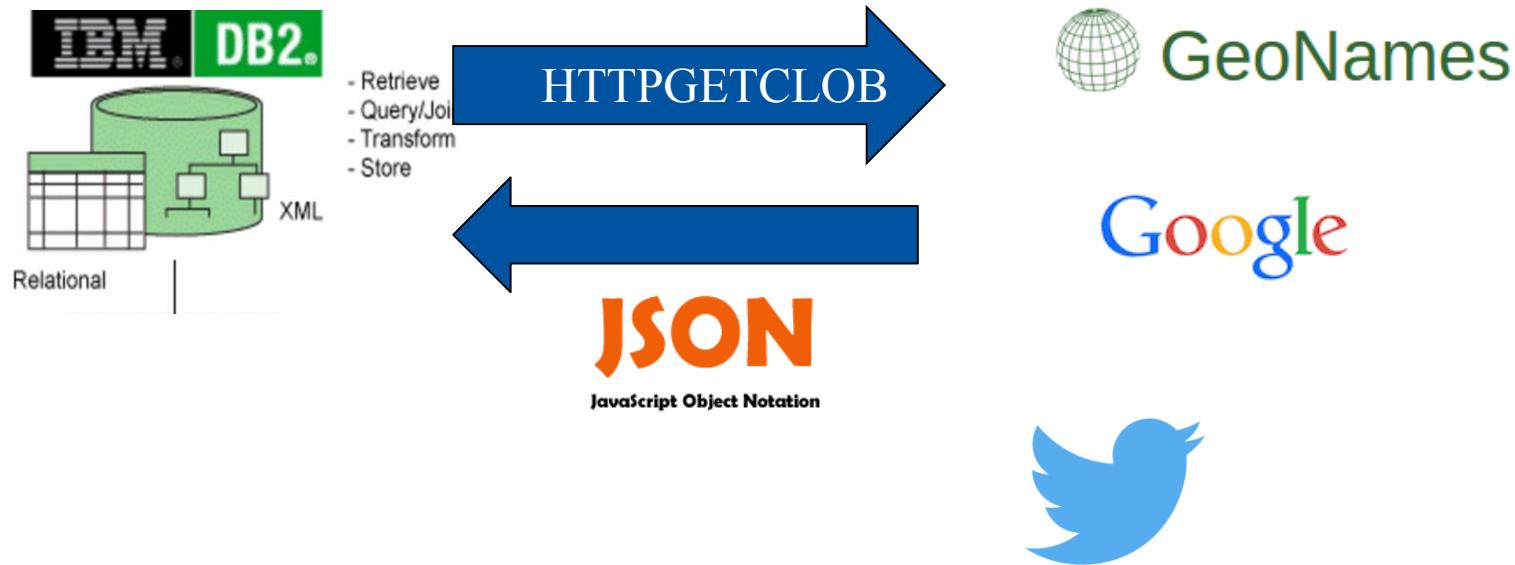
@partNum	productName	quantity	USPrice	shipDate
872-AA	Lawnmower	1	149.99	2014-11-20
945-ZG	Sapphire Bracelet	2	178.99	<null>

2 record(s) selected

Require
PI39003

Db2 for z/OS as REST service consumer

YAHOO![®]



Db2 for z/OS as REST service consumer



```
SELECT DB2XML.HTTPGETCLOB(
    CAST ('http://ws.geonames.org/countryInfo?lang=' ||
          DB2XML.URLENCODE('en', '') ||
          '&country=' ||
          DB2XML.URLENCODE('us', '') ||
          '&type=JSON' AS VARCHAR(255)),
    CAST(NULL AS CLOB(1K)))
FROM SYSIBM.SYSDUMMY1
```

Output (in JSON)

```
{"status": {
    "message": "...",
    "value": 10}}
```

with tablea as (

```
SELECT DB2XML.HTTPGETCLOB(
    CAST ('http://ws.geonames.org/countryInfo?lang=' ||
          DB2XML.URLENCODE('en', '') ||
          '&country=' ||
          DB2XML.URLENCODE('us', '') ||
          '&type=JSON' AS VARCHAR(255)),
    CAST(NULL AS CLOB(1K))) as DATA
FROM SYSIBM.SYSDUMMY1)
```

```
select JSON_VAL(SYSTOOLS.JSON2BSON(DATA),
  'status.value', 'i')
from tablea
```

10

1 record(s) selected

Db2 for z/OS as REST service consumer - Yahoo REST APIs

YAHOO!



With TABLEA as (

```
SELECT DB2XML.HTTPGETCLOB( CAST
('https://query.yahooapis.com/v1/public/yql?q=select%20*%20from%20weather.forecast%20where%20woeid%20in%20%28select%20woeid%20from%20geo.places%281%29%20where%20text%3D%22fremont%2C%20ca%22%29&format=json&env=store%3A%2F%2Fdatatables.org%2Falltableswithkeys' AS VARCHAR(255)),
CAST(NULL AS CLOB(1K))) AS DATA
FROM SYSIBM.SYSDUMMY1)
select
  JSON_VAL(SYSTOOLS.JSON2BSON(X.VALUE), 'date', 's:20') as date,
  JSON_VAL(SYSTOOLS.JSON2BSON(X.VALUE), 'high', 'i') as temperature_high,
  JSON_VAL(SYSTOOLS.JSON2BSON(X.VALUE), 'low', 'i') as temperature_low,
  JSON_VAL(SYSTOOLS.JSON2BSON(X.VALUE), 'text', 's:20') as description
from tableau, TABLE(SYSTOOLS.JSON_TABLE(SYSTOOLS.JSON2BSON(DATA),
'query.results.item.forecast', 's:200')) X
```

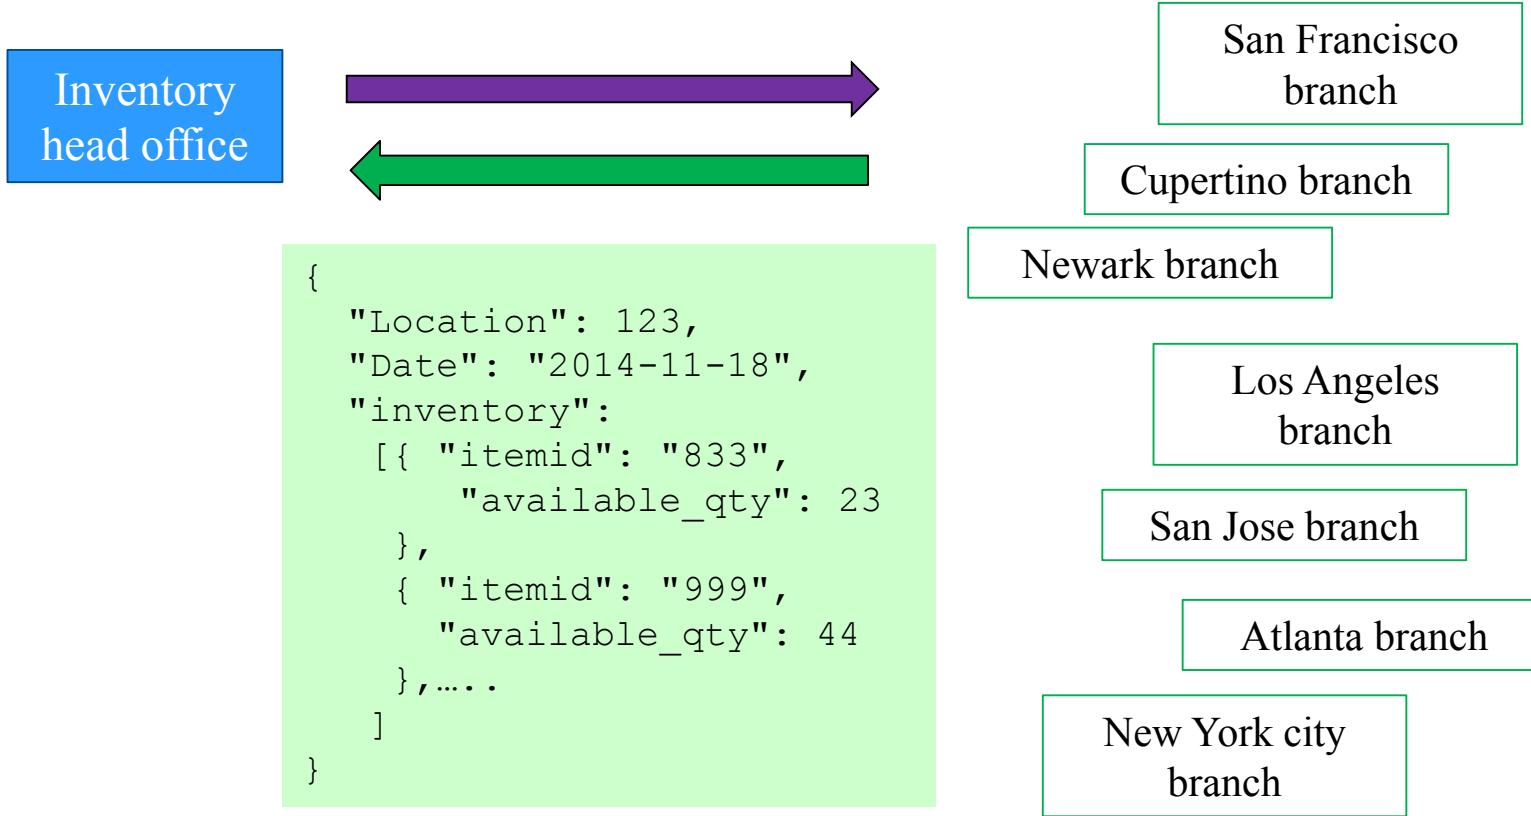
Array output, use
JSON_TABLE

DATE	TEMPERATURE_HIGH	TEMPERATURE_LOW	DESCRIPTION
06 Aug 2016	75	54	Partly Cloudy
07 Aug 2016	76	53	Mostly Sunny
08 Aug 2016	78	53	Mostly Sunny
09 Aug 2016	80	56	Mostly Sunny
10 Aug 2016	79	57	Mostly Sunny
11 Aug 2016	77	56	Mostly Sunny

...

10 record(s) selected

Db2 for z/OS as REST service consumer – inventory count



Db2 for z/OS as REST service consumer – inventory count

Inventory
head office

```
with tablea as (
  SELECT DB2XML.HTTPGETCLOB(
    'http://....', ...
  ) as DATA
  FROM SYSIBM.SYSDUMMY1)
  SELECT
    JSON_VAL(SYSTOOLS.JSON2BSON(X.VALUE), 'itemid', 's:10') as
      "itemid",
    JSON_VAL(SYSTOOLS.JSON2BSON(X.VALUE), 'available_qty', 's:20') as
      "available_qty"
  FROM tablea,
    TABLE(SYSTOOLS.JSON_TABLE(SYSTOOLS.JSON2BSON(DATA), 'inventory',
    's:200')) X
```

```
{
  "Location": 123,
  "Date": "2014-11-18",
  "inventory":
    [ { "itemid": "833",
        "available_qty": 23
      },
      { "itemid": "999",
        "available_qty": 44
      }
    ]
}
```

itemid	available_qty
833	23
999	44
2 record(s) selected	

Db2 for z/OS as REST consumer – inventory count - JOIN

ITEMTABLE

PART	DESCRIPTION
833	Lawnmower
999	Sapphire Bracelet

```

with tablea as (
  SELECT DB2XML.HTTPGETCLOB(
    'http://....', ...
  ) as DATA
  FROM SYSIBM.SYSDUMMY1),
  breakdown (itemid, available_qty) as
  (
    SELECT
      JSON_VAL(SYSTOOLS.JSON2BSON(X.VALUE), 'itemid', 's:10') as "itemid",
      JSON_VAL(SYSTOOLS.JSON2BSON(X.VALUE), 'available_qty', 's:20') as
    "available_qty"
    FROM tablea,
      TABLE(SYSTOOLS.JSON_TABLE(SYSTOOLS.JSON2BSON(DATA), 'inventory',
    's:200')) X)
  SELECT B.itemid, B.available_qty, I.description
  FROM ITEMTABLE I, breakdown B
  WHERE B.itemid = I.part

```

JOIN

```
{
  "Location": 123,
  "Date": "2014-11-18",
  "inventory":
  [ { "itemid": "833",
    "available_qty": 23
  },
  { "itemid": "999",
    "available_qty": 44
  }, ...
]
```

itemid	available_qty
833	23
999	44

2 record(s) selected

Db2 for z/OS as REST service consumer – inventory count - JOIN

ITEMTABLE

PART	DESCRIPTION
833	Lawnmower
999	Sapphire Bracelet

From ITEMTABLE

itemid	available_qty
833	23
999	44

2 record(s) selected

From HTTPGETCLOB

ITEMID	AVAILABLE_QTY	DESCRIPTION
833	23	Lawnmower
999	44	Sapphire Bracelet

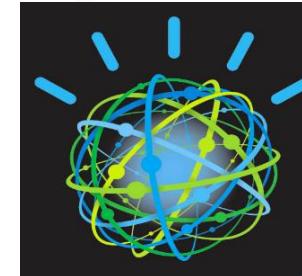
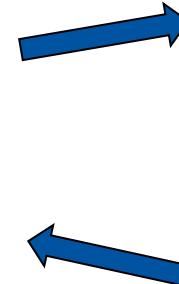
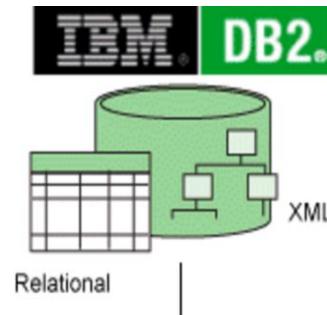
2 record(s) selected

RESULT

Agenda

- Db2 as REST provider using Db2 native REST service
- Db2 as REST consumer using JSON feature in Db2
- **Cognitive Db2 solution using Waston**
- Perform Machine Learning on Db2 data without a GUI
- Spark analytics

Db2 for z/OS with Watson



```
{
  "document_tone": {
    "tone_categories": [
      {
        "tones": [
          {
            "score": 0.18671406839712007,
            "tone_id": "Anger",
            "tone_name": "Anger",
            "tone_category_id": "emotion_tone",
            "tone_category_name": "Emotion Tone"
          }
        ]
      }
    ]
  }
}
```

- View a JSON analysis of your content

Results

Emotion (probability score)

Anger		.95	LIKELY
Disgust		.40	UNLIKELY
Fear		.20	UNLIKELY
Joy		.02	UNLIKELY
Sadness		.13	

Writing Style (raw score)

Analytical		.92
Confident		0.0
Tentative		0.0

Social Tendencies (raw score)

Agreeableness		.92
Conscientiousness		.20
Emotional Range		.90
Extraversion		.70
Openness		.07

<http://www.ibm.com/watson/developercloud/doc/tone-analyzer/case-study.shtml>

What tone is it?

I am so happy today!



I have worked days and nights
for three years without a salary
raise!

I thought I can sleep in this session, but that
crazy speaker keep asking us questions. Now
she is asking us to read this long text and tell
her what tone it is. How do I know? Ask a
machine to do this....

???

Db2 for z/OS and Watson Tone Analyzer

```

WITH Authorization AS (
    SELECT DB2XML.BASE64ENCODE(CAST('<userid>:<password>' AS
VARCHAR(64) CCSID 1208)) AS Authorization_Identity
FROM SYSIBM.SYSDUMMY1
)
SELECT DB2XML.HTTPGETCLOB(
    CAST ('https://gateway.watsonplatform.net/tone-
analyzer/api/v3/tone?version=2016-05-19&text=' ||
DB2XML.URLENCODE('I thought I can sleep in this session, but that
crazy speaker keep asking us questions. Now she is asking us to
read this long text and tell her what tone it is. How do I know?
Ask a machine to do this','') AS VARCHAR(255)),
    CAST('<httpHeader><header name="Authorization" value="Basic
'||Authorization_Identity||'" /></httpHeader>' AS CLOB(1K)))
FROM Authorization#

```

```
{"document_tone": {"tone_categories": [{"tones": [
{"score": 0.206372, "tone_id": "anger", "tone_name": "Anger"}, 
{"score": 0.364932, "tone_id": "disgust", "tone_name": "Disgust"}, 
 {"score": 0.451429, "tone_id": "fear", "tone_name": "Fear"}, 
 {"score": 0.097665, "tone_id": "joy", "tone_name": "Joy"}, 
 {"score": 0.439061, "tone_id": "sadness", "tone_name": "Sadness"}]}}
```

Hmm...What
JSON function
should I used?

Db2 for z/OS and Waston Tone Analyzer

```

WITH Authorization AS (
  SELECT DB2XML.BASE64ENCODE(CAST('<userid>:<password>' AS VARCHAR(64) CCSID
1208)) AS Authorization_Identity
  FROM SYSIBM.SYSDUMMY1
),
TONETABLE AS (
  SELECT DB2XML.HTTPGETCLOB(
    CAST ('https://gateway.watsonplatform.net/tone-
analyzer/api/v3/tone?version=2016-05-19&text=' || DB2XML.URLENCODE('I thought
I can sleep in this session, but that crazy speaker keep asking us questions.
Now she is asking us to read this long text and tell her what tone it is. How
do I know? Ask a machine to do this','') AS VARCHAR(255)),
    CAST('<httpHeader><header name="Authorization" value="Basic
'||Authorization_Identity||'" /></httpHeader>' AS CLOB(1K))) TONERESULT
  FROM Authorization)
select
  JSON_VAL(SYSTOOLS.JSON2BSON(X.VALUE), 'score', 'f') as score,
  JSON_VAL(SYSTOOLS.JSON2BSON(X.VALUE), 'tone_name', 's:20') as tone
from TONETABLE,
TABLE(SYSTOOLS.JSON_TABLE(SYSTOOLS.JSON2BSON(TONERESULT),
'document_tone.tone_categories.tones', 's:200')) X
order by score DESC

```

Db2 for z/OS and Waston Tone Analyzer

```

WITH Authorization AS (
  SELECT DB2XML.BASE64ENCODE
        1208) ) AS Authorization_Identity
  FROM SYSIBM.SYSDUMMY1
),
TONETABLE AS (
  SELECT DB2XML.HTTPGETCLOB(
    CAST ('https://gateway.watsonplatform.net/tone-analyzer/api/v3/tone?version=2017-05-26' AS CLOB(1K)), TONERESULT
    I can sleep in this session, 0.007
    Now she is asking us to read 0.004
    do I know? Ask a machine to 0.0
    CAST('<httpHeader><head>
  ||Authorization_Identity||'
  FROM Authorization)
select
  JSON_VAL(SYSTOOLS.JSON2BSON(X.VALUE), 'score', 'f') as score,
  JSON_VAL(SYSTOOLS.JSON2BSON(X.VALUE), 'tone_name', 's:20') as tone
from TONETABLE,
TABLE(SYSTOOLS.JSON_TABLE(SYSTOOLS.JSON2BSON(TONERESULT),
'document_tone.tone_categories.tones', 's:200')) X
order by score DESC

```

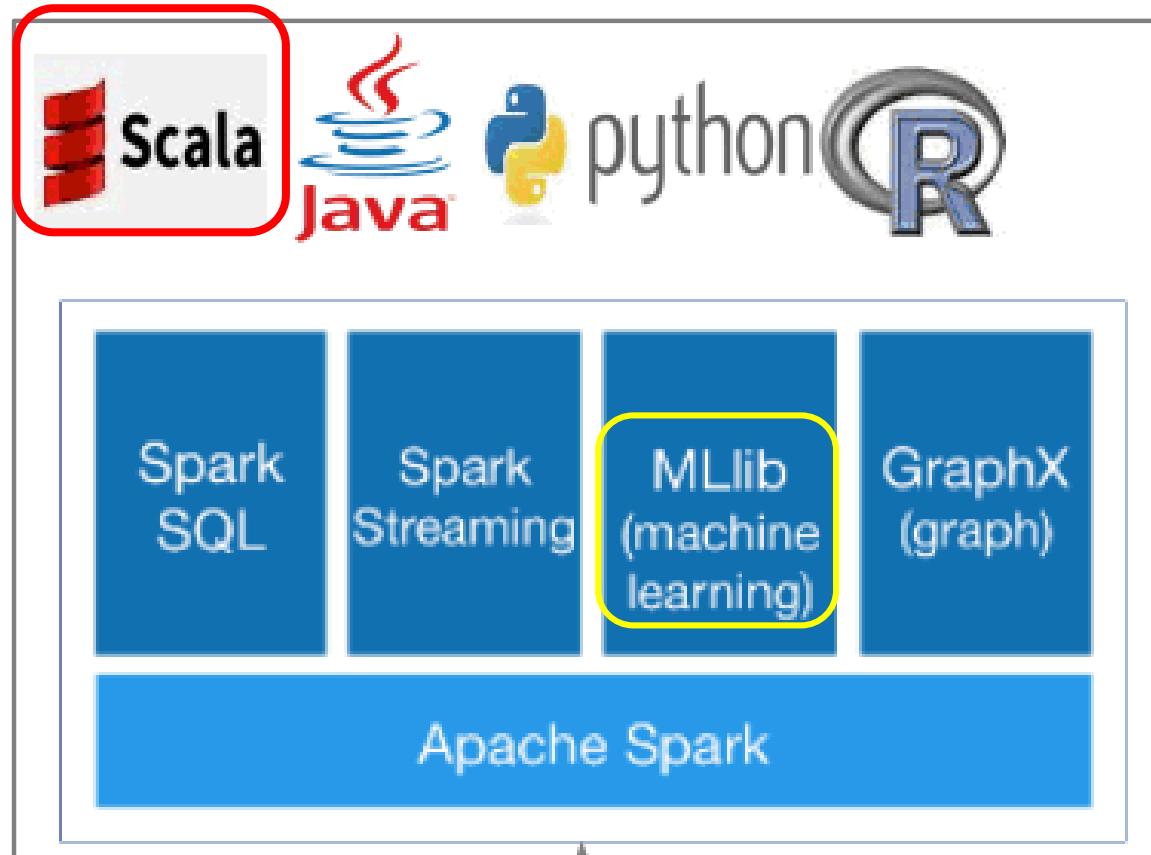
SCORE	TONE
0.99	Extraversion
0.987	Emotional Range
0.97	Agreeableness
0.97	Analytical
0.699	Tentative
0.451429	Fear
0.439061	Sadness
0.364932	Disgust
0.206372	Anger
0.097665	Joy
0.007	Openness
0.004	Conscientiousness
0.0	Confident

13 record(s) selected

Agenda

- Db2 as REST provider using Db2 native REST service
- Db2 as REST consumer using JSON feature in Db2
- Cognitive Db2 solution using Waston
- Perform Machine Learning on Db2 data without a GUI
- Spark analytics

Spark Overview



Installation/Pre-req

1. Db2 JDBC driver jar :
 db2jcc4.jar,
 db2jcc_license_cisuz.jar
 (the license file is only
 needed if the Db2 for
 z/OS server is not
 activated)
2. Java 1.7 or above
3. Spark (I use Spark 2.1.0).
 You can download the
 pre-built version from
[https://spark.apache.org/
 downloads.html](https://spark.apache.org/downloads.html)

Our Scenarios

AGE	GENDER	BP_TOP	CHOLESTEROL	SODIUM	POTASSIUM	DRUG
			_RATIO			
23	F	139.0	5.1	0.793	0.031	drugY
47	M	90.0	5.5	0.739	0.056	drugC
47	M	90.0	4.9	0.697	0.069	drugC
28	F	120.0	4.8	0.564	0.072	drugX
61	F	90.0	5.2	0.559	0.031	drugY
22	F	115.0	4.3	0.677	0.079	drugX
49	F	119.0	4.3	0.790	0.049	drugY
41	M	90.0	4.8	0.767	0.069	drugC
60	M	120.0	4.7	0.777	0.051	drugY
43	M	90.0	2.2	0.526	0.027	drugY

Input (features)



Output

All string need to be converted to double!

Steps Overview

1. Create and populate table
2. Launch Spark and connect to Db2
3. Preprocess the inputs(features) to our model
4. Find a suitable machine learning algorithms
5. Train and test the model
6. Save the model

For Step 2 to 6, we are following the steps in
<https://spark.apache.org/docs/2.0.2/ml-classification-regression.html#decision-tree-classifier>

Step 1: Create and Populate table

```
drop table t1#
create table t1(age int, gender char(1), BP_top double,
cholesterol_ratio double, sodium double, potassium double, drug
varchar(5))#
insert into t1 values(23, 'F', 139, 5.1, 0.793, 0.031, 'drugY')#
insert into t1 values(47, 'M', 90, 5.5, 0.739, 0.056, 'drugC')#
insert into t1 values(47, 'M', 90, 4.9, 0.697, 0.069, 'drugC')#
insert into t1 values(28, 'F', 120, 4.8, 0.564, 0.072, 'drugX')#
insert into t1 values(61, 'F', 90, 5.2, 0.559, 0.031, 'drugY')#
insert into t1 values(22, 'F', 115, 4.3, 0.677, 0.079, 'drugX')#
insert into t1 values(49, 'F', 119, 4.3, 0.790, 0.049, 'drugY')#
insert into t1 values(41, 'M', 90, 4.8, 0.767, 0.069, 'drugC')#
insert into t1 values(60, 'M', 120, 4.7, 0.777, 0.051, 'drugY')#
insert into t1 values(43, 'M', 90, 2.2, 0.526, 0.027, 'drugY')#
```

Step 2: Launch Spark and connect to Db2 for z/OS

```
Set HADOOP_HOME=C:\Hadoop  
set SPARKBIN=C:\Spark210\spark-2.1.0-bin-hadoop2.7\bin  
SET PATH=C:\Program Files\IBM\Java70\bin;%SPARKBIN%;%PATH%
```

Run as
Administrator

```
C:\Spark210\spark-2.1.0-bin-hadoop2.7\bin\spark-shell.cmd --driver-class-path  
C:\jcc_home\jccbuilds\jcc411\db2jcc4.jar;C:\jcc_home\jccbuilds\jcc411\db2jc  
c_license_cisuz.jar
```

```
import org.apache.spark.ml.classification.DecisionTreeClassifier  
import org.apache.spark.ml.feature.{IndexToString, StringIndexer,  
VectorIndexer, VectorAssembler, OneHotEncoder, SQLTransformer}  
import org.apache.spark.ml.{Pipeline, PipelineModel}  
import org.apache.spark.ml.evaluation.MulticlassClassificationEvaluator  
import org.apache.spark.ml.classification.DecisionTreeClassificationModel  
import org.apache.spark.sql.SparkSession
```

```
//Create a Spark session (Since Spark 2.0, use SparkSession instead of sqlContext)
```

```
val spark =  
SparkSession.builder().appName("JaneMLEExample").getOrCreate()
```

Step 2: Launch Spark and connect to Db2 for z/OS

```
//Get data from Db2 (customize according to your setting)
val DB2Data = spark.read.format("jdbc").options(Map("url" ->
"jdbc:db2://hostname:port/location:currentSchema=schema;user=user;passwo
rd=password;","driver" -> "com.ibm.db2.jcc.DB2Driver", "dbtable" ->
"SYSADM.T1")).load()
```

Step 3: Preprocess the inputs(features) to our model

```
//Convert all string to double
val labelIndexer = new
StringIndexer().setInputCol("DRUG").setOutputCol("drugLabel").fit(DB2Data)
val genderIndexer = new
StringIndexer().setInputCol("GENDER").setOutputCol("genderIndex")
val genderEncoder = new
OneHotEncoder().setInputCol("genderIndex").setOutputCol("genderVec")
```

Step 3: Preprocess the inputs(features) to our model – cont'd

```
//Create features vector as input to our model
val assembler = new VectorAssembler().setInputCols(Array("AGE", "genderVec",
"BP_TOP", "CHOLESTEROL_RATIO", "SODIUM", "POTASSIUM")).setOutputCol("features")

//Split the data into training and test sets (30% held out for testing).
val Array(trainingData, testData) = DB2Data.randomSplit(Array(0.7, 0.3))
```

Step 4: Find a suitable machine learning algorithms

```
//Train a DecisionTree model.
val dt = new
DecisionTreeClassifier().setLabelCol("drugLabel").setFeaturesCol("features")

//Convert indexed label (prediction in double) back to original labels like drugX, drugY, drugC (output in predictedLabel)
val labelConverter = new
IndexToString().setInputCol("prediction").setOutputCol("predictedLabel").
setLabels(labelIndexer.labels)
```

Step 5: Train and test the model

```
//Chain all the steps (like indexers, tree) in a Pipeline
```

```
val pipeline = new Pipeline().setStages(Array(labelIndexer,  
genderIndexer, genderEncoder, assembler, dt, labelConverter))
```

```
//Train model using trainingData
```

```
val model = pipeline.fit(trainingData)
```

```
//Make predictions using testData .
```

```
val predictions = model.transform(testData)
```

```
//Display predication  
predictions.select("DRUG",  
"predictedLabel",  
"drugLabel",  
"features").show()
```

Actual	Predicted	drugLabel	features
DRUG	predictedLabel		
drugX		drugX	2.0 [22.0,0.0,115.0,4....
drugC		drugC	1.0 [47.0,1.0,90.0,4....
drugY		drugY	0.0 [49.0,0.0,119.0,4....

Step 6: Save the model

```
//save the fitted pipeline to disk
model.write.overwrite().save("C:/Jane/Work/DB2/socialMedia/machineLearning/DB2Model1")
```

//load it back in during production

```
val sameModel =
PipelineModel.load("C:/Jane/Work/DB2/socialMedia/machineLearning/DB2Model1")
```



```
val predictionAgain = sameModel.transform(testData)
```

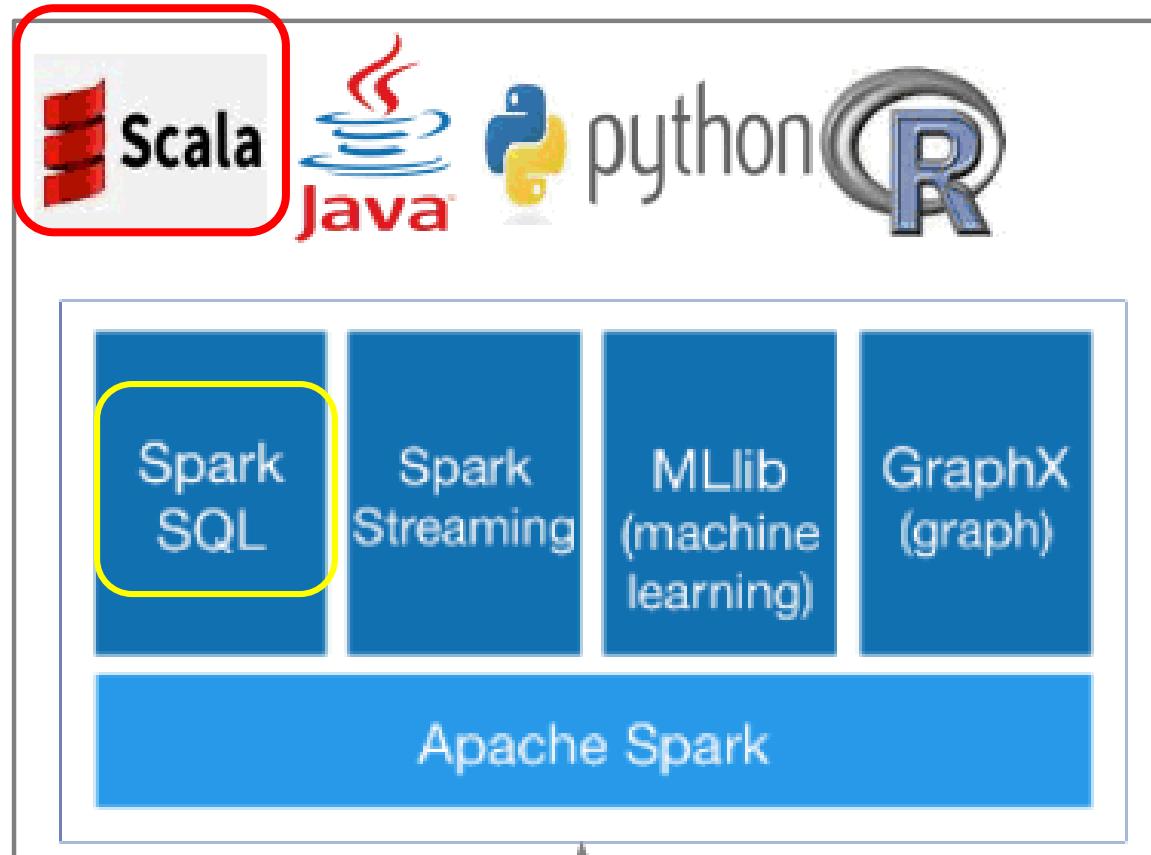


Spark 1.6.2 will
NOT work

Agenda

- Db2 as REST provider using Db2 native REST service
- Db2 as REST consumer using JSON feature in Db2
- Cognitive Db2 solution using Waston
- Perform Machine Learning on Db2 data without a GUI
- Spark analytics

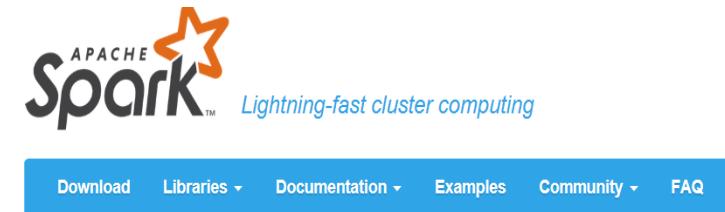
Spark Overview



Installation/Pre-req

1. JDK 1.7 or 1.8
2. Download spark* from
<https://spark.apache.org/downloads.html>
(I download the pre-built binaries to my Window 7 machine)
3. JCC jars

Ref: <http://www.worldofdb2.com/profiles/blogs/using-spark-s-interactive-scala-shell-for-accessing-db2-data>



Download Apache Spark™

Our latest stable version is Apache Spark 1.6.2, released on June 25, 2016 ([release notes](#)) ([git tag](#))

1. Choose a Spark release:
2. Choose a package type:
3. Choose a download type:
4. Download Spark: [spark-1.6.2-bin-hadoop2.6.tgz](#)
5. Verify this release using the [1.6.2 signatures and checksums](#) and [project release KEYS](#).

Note: Scala 2.11 users should download the Spark source package and build with Scala 2.11 support.

Running Scala : Environment Setting

1. :: MUST start cmd prompt as administrator!!

I put this is
a .bat file

```
Set HADOOP_HOME=C:\Hadoop
set SPARKBIN=C:\Spark162\spark-1.6.2-bin-hadoop2.6\bin
::add JCC to classpath
set JCCJAR=C:\jcc_home\jccbuids\jcc411
set
CLASSPATH=%JCCJAR%\db2jcc4.jar;%JCCJAR%\db2jcc_license_ci
suz.jar;%JCCJAR%\db2jcc_license_cu.jar;%CLASSPATH%
::set path for jdk 7
SET PATH=C:\Program
Files\IBM\Java70\bin;%SPARKBIN%;%PATH%
```

2. Run spark-shell.cmd under bin directory

```
C:\Spark162\spark-1.6.2-bin-hadoop2.6\bin\spark-shell.cmd
```

Basic Scala

3. Create a dataframe to a table inside Db2 for z/OS

When a SQL context is created, you should see a scala prompt:



```
a0134ed-4f5f-49ee-97a4-5182af9c7374/_tmp_space.db
16/07/16 18:41:10 INFO SparkILoop: Created sql context (with Hive support)..
SQL context available as sqlContext.

scala> _
```

For example, the command to create a dataframe to SYSIBM. SYSXMLSTRINGS catalog table:

```
val catalogDF = sqlContext.load("jdbc", Map("url" ->
"jdbc:db2://dtec222.vmecl.svl.ibm.com:446/STLEC1:currentSchema=s
ysadm;user=sysadm;password=password;","driver" ->
"com.ibm.db2.jcc.DB2Driver","dbtable" ->
"SYSIBM.SYSXMLSTRINGS"))
```

Basic Scala Command

- Enter following scala command to show the content of dataframe (SYSIBM. SYSXMLSTRINGS)

```
catalogDF.show()
```

```
catalogDF.registerTempTable("xmldf")
```

```
sqlContext.sql("select STRINGID, STRING from xmldf
WHERE STRINGID < 1005").show
```

```
+-----+-----+
|STRINGID|      STRING|
+-----+-----+
|    1001|  product|
|    1002|description|
|    1003|     name|
|    1004|    detail|
+-----+-----+
```

STRINGID	STRING	IBMREQDI
1001	product	NI
1002	description	NI
1003	name	NI
1004	detail	NI
1005	http://www.w3.org...	NI
1006	space	NI
1007	http://posample.org/	NI
1008	pid	NI
1009	details	NI
1010	price	NI
1011	weight	NI
1021	a	NI
1022	b	NI
1023	c	NI
1024	table2	NI
1025	request	NI
1026	should	NI
1027	id	NI
1028	noAggl	NI
1029	requestid	NI

Spark Analytics on XML

You can use either
one or both

2 major packages:

1. `scala.xml.XML`

(already include in the pre-built binaries, no need to download extra package)

Good for basic
searching

2. `com.databricks.spark.xml`

Use the following command to pull in the databricks XML package when
launching the scala shell:

```
bin\spark-shell --packages com.databricks:spark-xml_2.10:0.3.3
```

* this command is for scala 2.10

See my blogs in Resources
section for more details

Good for XML document
with **array structure**

Our Scenarios

```
create table t1(id int, xmlcol xml);
```

2 XML documents are inserted in T1

```
<PO id="123" orderDate="2013-11-18">
  <customer cid="123"/>
  <items>
    <item partNum="872-AA">
      <productName>Lawnmower</productName>
      <quantity>1</quantity>
      <USPrice>149.99</USPrice>
      <shipDate>2013-11-20</shipDate>
    </item>
    <item partNum="945-ZG">
      <productName>Sapphire
      Bracelet</productName>
      <quantity>2</quantity>
      <USPrice>178.99</USPrice>
      <comment>Not shipped</comment>
    </item>
  </items>
</PO>
```

Hubert

```
<PO id="999" orderDate="2013-11-18">
  <customer cid="999"/>
  <items>
    <item partNum="900-AA">
      <productName>pen</productName>
      <quantity>1</quantity>
      <USPrice>5.0</USPrice>
      <shipDate>2013-11-20</shipDate>
    </item>
    <item partNum="145-ZG">
      <productName>Johnson lotion</productName>
      <quantity>2</quantity>
      <USPrice>10.0</USPrice>
      <comment>Not shipped</comment>
    </item>
    <item partNum="945-ZG">
      <productName>Sapphire Bracelet</productName>
      <quantity>3</quantity>
      <USPrice>178.99</USPrice>
      <comment>Not shipped</comment>
    </item>
  </items>
</PO>
```

Jorn

scala.xml.XML

//1. import the required package

```
scala> import scala.xml.XML
```

//2. create a dataframe for our T1 table

```
scala> val xmlDF = sqlContext.load("jdbc", Map("url" ->
"jdbc:db2://dtec222.vmec.svl.ibm.com:446/STLEC1:currentSchema=sysadm;user=sysadm;password=password;", "driver" ->
"com.ibm.db2.jcc.DB2Driver", "dbtable" -> "SYSADM.T1"))
```

//3. create a string array to hold the XML document

```
scala> var xmlStringArray = (xmlDF.select(xmlDF("XMLCOL"))).map(
_.getString( 0 ) ).map( scala.xml.XML.loadString( _ ) ).collect()
```

```
xmlStringArray: Array[scala.xml.Elem] = Array(<PO
orderDate="2013-11-18" id="123"><customer cid="123"/><items>
<item partNum="872-AA"><productName>Lawnmower
</productName><quantity>1</quantity><USPrice>149.99</USPrice><sh
ipDate>2013-11-20</shipDate>.....
```

<PO ..id="123">...

<PO ..id="999">...

scala.xml.XML

//find all cid

```
scala> var cid = xmlStringArray.map( _ \\ "PO" \ "customer" \ "@cid" )
cid: Array[scala.xml.NodeSeq] = Array(123, 999)
```

//create a function to check for a particular cid

```
scala> def cidEquals(value: String)(node: scala.xml.Node) = {
    ( node \\ "PO" \ "customer" \ "@cid" ).exists(_.text == value)
}
```

//find all productName for cid = 123

```
scala> xmlStringArray.map( _ \\ "_" filter cidEquals("123") ).map(
    _ \\ "PO" \ "items" \ "item" \ "productName" ).foreach(println(_))
```

```
<productName>Lawnmower</productName><productName>Sapphire
Bracelet</productName>
```

```
<PO id="123" orderDate="2013-11-18"
  <customer cid="123"/>
  <items>
    <item partNum="872-AA">
      <productName>Lawnmower</productName>
    ....
```

```
<PO id="999" orderDate="2013-11-18">
  <customer cid="999"/>
  <items>
    <item partNum="900-AA">
      <productName>pen</productName>
    ....
```

Hubert: What
happen to my order?

scala.xml.XML

Hubert: How
much I need to
pay?

//create a function to find total price

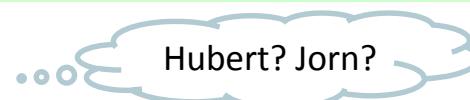
```
scala> def totalPrice(node: scala.xml.Node) :  
    Double = {  
        var USPrice :Double = 0  
        var quantity :Int = 0  
        var totalPrice :Double = 0  
        var items = (node \\ "PO" \\ "items")  
        for (item <- items \\ "item")  
        {  
            USPrice = (item \\ "USPrice").text.toDouble  
            quantity = (item \\ "quantity").text.toInt  
            totalPrice = totalPrice + (USPrice * quantity)  
        }  
  
        return totalPrice  
    }
```

```
<PO id="123" orderDate="2013-11-18">  
  <customer cid="123"/>  
  <items>  
    <item partNum="872-AA">  
      <productName>Lawnmower</productName>  
      <quantity>1</quantity>  
      <USPrice>149.99</USPrice>  
      <shipDate>2013-11-20</shipDate>  
    </item>  
    <item partNum="945-ZG">  
      <productName>Sapphire  
      Bracelet</productName>  
      <quantity>2</quantity>  
      <USPrice>178.99</USPrice>  
      <comment>Not shipped</comment>  
    </item>  
  </items>  
</PO>
```

scala.xml.XML

//find total price for each XML document

```
scala> xmlStringArray.map(totalPrice(_)).foreach(println(_))
507.97
561.97
```



//associate cid with total price

```
scala> def cidNTotalPriceSimple(node: scala.xml.Node) = {
  ((node \\ "PO" \ "customer" \ "@cid").text.toInt, totalPrice(node))
}
```

//create a dataframe for cid and corresponding total price

```
scala> val df =
sc.parallelize(xmlStringArray.map(cidNTotalPriceSimple(_))).toDF("c
id", "totalPrice")
```

```
scala> df.printSchema()
root
```

```
|-- cid: integer (nullable = false)
|-- totalPrice: double (nullable = false)
```

```
<PO id="999" orderDate="2013-11-18">
<customer cid="999"/>
<items>....
```

cid	totalPrice
123	507.97
999	561.97

scala.xml.XML

//find all cid and total price and sort total price ascendingly

```
scala> df.select("cid", "totalPrice").sort(asc("totalPrice")).show
+---+-----+
|cid|totalPrice|
+---+-----+
|123|      507.97|
|999|      561.97|
+---+-----+
```

//find cid with total price > 550

```
df.filter($"totalPrice" > 550).show
+---+-----+
|cid|totalPrice|
+---+-----+
|999|      561.97|
+---+-----+
```

How about
attributes w/
prefix?

```
scala> val xmldoc = <a
  xmlns:bar="http://bar.com"
  bar:b="bAttributes"></a>
```

```
scala> xmldoc \
  "@{http://bar.com}b"
```

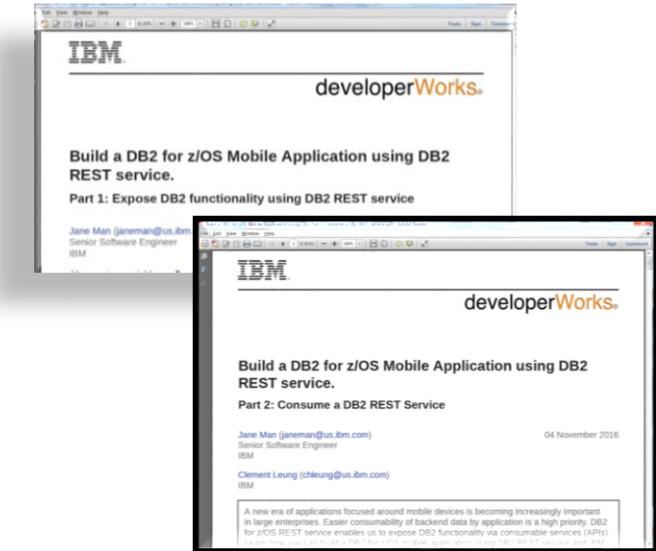
```
res0: scala.xml.NodeSeq =
bAttributes
```

Summary

- **Db2 + Watson:** Cognitive Solution on your enterprise data
- Use Db2 for z/OS as **REST consumer** to utilize data outside Db2
- Db2 for z/OS is a **REST service provider** now!
- Apache **Spark** provides big data analytics capabilities on Db2 data.
- Write your own **Machine Learning** application using Db2 for z/OS data

Resources

- Db2 for z/OS Native DDF REST services - 2 whitepaper by Jane Man
https://www.ibm.com/developerworks/community/blogs/e429a8a2-b27f-48f3-aa73-ca13d5b69759/entry/DB2_for_z_OS_Native_DDF_REST_services?lang=en
- Db2 REST services
http://www.ibm.com/support/knowledgecenter/en/SSEPEK_11.0.0/restserv/src/tpc/db2z_restservices.html
- IBM MobileFirst 7 Developer Edition
<https://developer.ibm.com/mobilefirstplatform/install/#studio>



Resources – Spark Analytics for XML

- IBM Packages for Apache Spark
<https://www.ibm.com/developerworks/java/jdk/spark/>
- Installing IBM z/OS Platform for Apache Spark
<http://www-03.ibm.com/support/techdocs/atsmastr.nsf/WebIndex/WP102609>
- Spark Analytics for XML data in DB2 for z/OS – Part 1 : scala.xml.XML by Jane Man
<http://ibm.co/2frZw2y>
- Spark Analytics for XML data in DB2 for z/OS – Part 2 : databricks XML package – on a single XML document by Jane Man <http://ibm.co/2frZE22>
- Spark Analytics for XML data in DB2 for z/OS – Part 3 : databricks XML package – on XML data in database by Jane Man <http://ibm.co/2frZF6a>
- Spark Blog 1 - Using Spark's interactive Scala shell for accessing DB2 data using JDBC driver and Spark's new DataFrames API
<http://www.worldofdb2.com/profiles/blogs/using-spark-s-interactive-scala-shell-for-accessing-db2-data>
- Instantiating a case class from a list of parameters <http://stackoverflow.com/questions/4290955/instantiating-a-case-class-from-a-list-of-parameters>

Resources – for Machine Learning and Python

- Machine Learning using DB2 for z/OS data and Spark Part 1 by Jane Man
https://www.ibm.com/developerworks/community/blogs/e429a8a2-b27f-48f3-aa73-ca13d5b69759/entry/Machine_Learning_using_DB2_for_z_OS_data_and_Spark_Part_1?lang=en
- Machine Learning using DB2 for z/OS data and Spark Part 2 by Jane Man
https://www.ibm.com/developerworks/community/blogs/e429a8a2-b27f-48f3-aa73-ca13d5b69759/entry/Machine_Learning_using_DB2_for_z_OS_data_and_Spark_Part_2?lang=en
- Quick Start in accessing DB2 for z/OS from Python applications by Jane Man
https://www.ibm.com/developerworks/community/blogs/e429a8a2-b27f-48f3-aa73-ca13d5b69759/entry/Quick_Start_in_accessing_DB2_for_z_OS_from_Python_applications?lang=en



Jane Man

IBM

janeman@us.ibm.com

Db2 for z/OS Modern Application: Watson, ML, REST, Spark and more

Session code: G12

*Please fill out your session
evaluation before leaving!*

